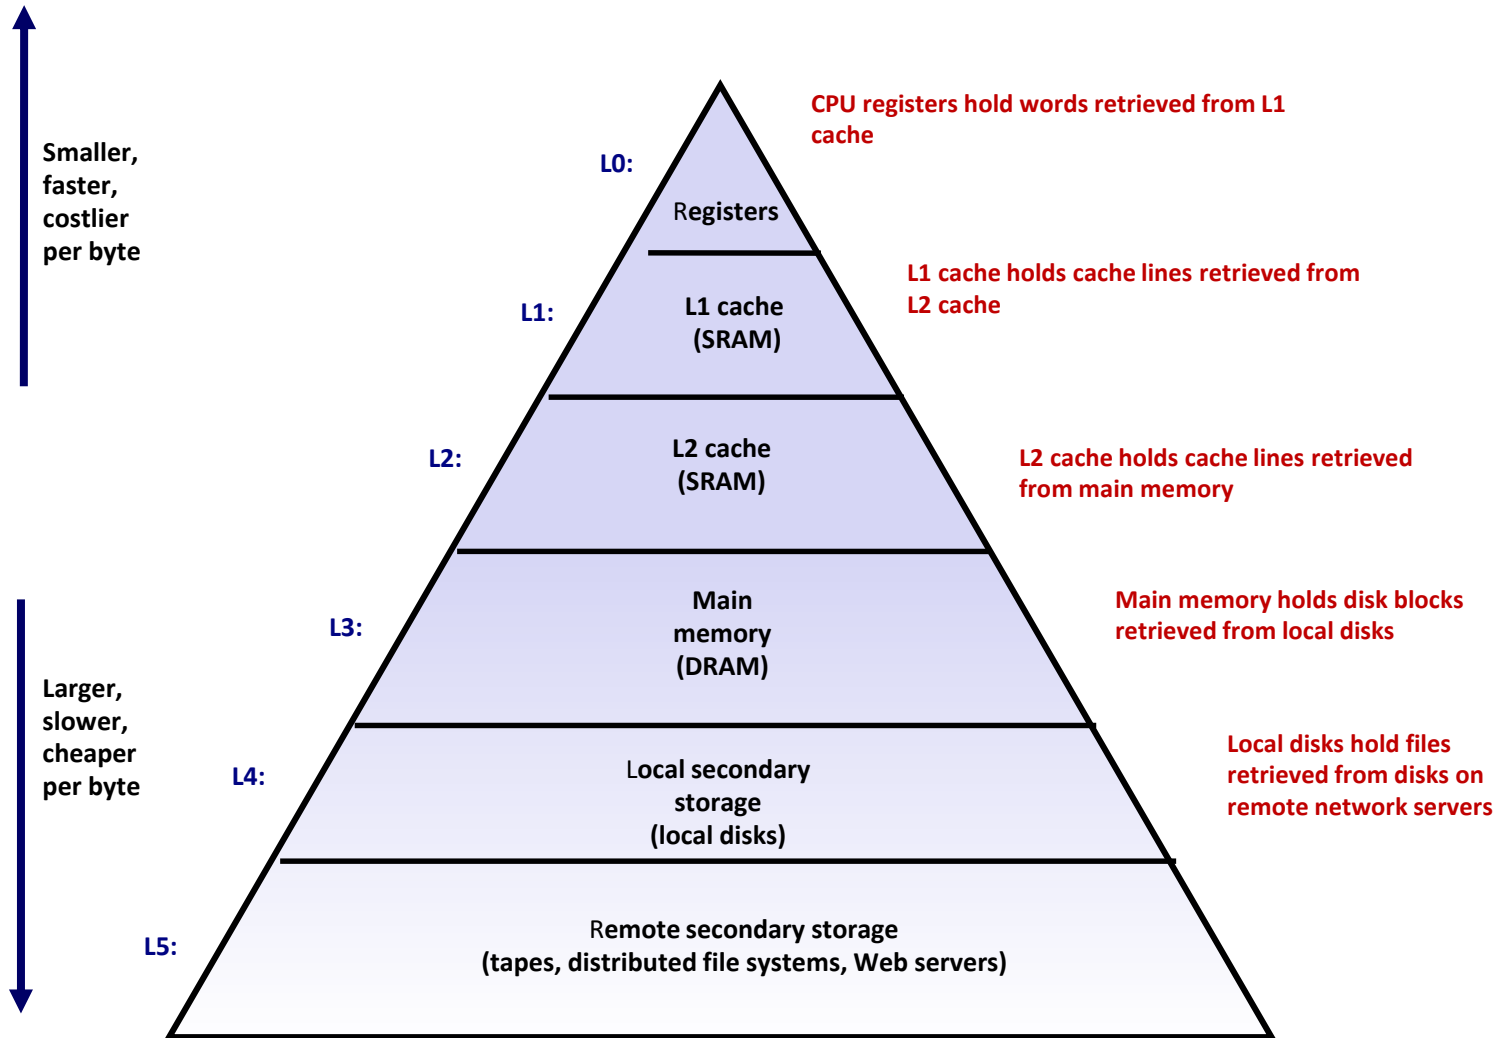


Cache Lab Implementation and Blocking

Outline

- **Memory organization**
- **Caching**
 - Different types of locality
 - Cache organization
- **Cache lab**
 - Cache Structure
 - getopt/fscanf/Malloc
- **Page Replacement**
 - LRU algorithm
 - FIFO algorithm

Memory Hierarchy



SRAM vs DRAM tradeoff

■ SRAM (cache)

- Faster (L1 cache: 1 CPU cycle)
- Smaller (Kilobytes (L1) or Megabytes (L2))
- More expensive and “energy-hungry”

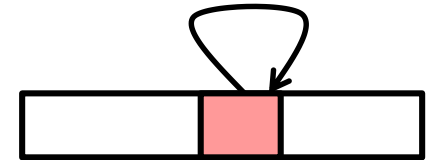
■ DRAM (main memory)

- Relatively slower (hundreds of CPU cycles)
- Larger (Gigabytes)
- Cheaper

Locality

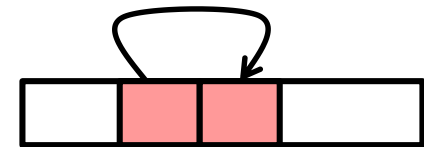
■ Temporal locality

- Recently referenced items are likely to be referenced again in the near future
- After accessing address X in memory, save the bytes in cache for future access



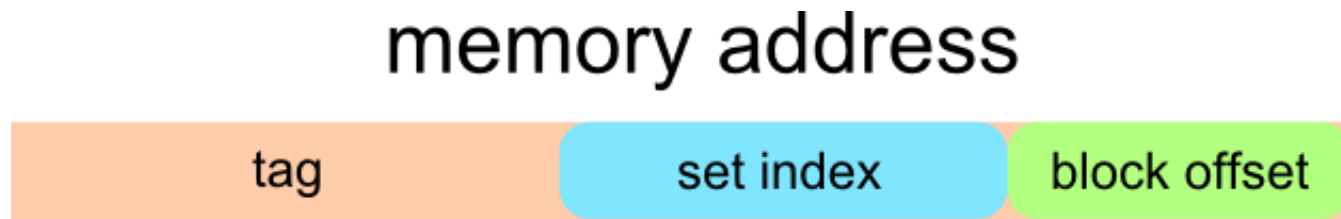
■ Spatial locality

- Items with nearby addresses tend to be referenced close together in time
- After accessing address X, save the block of memory around X in cache for future access



Memory Address

- For example, 64-bit on shark machines

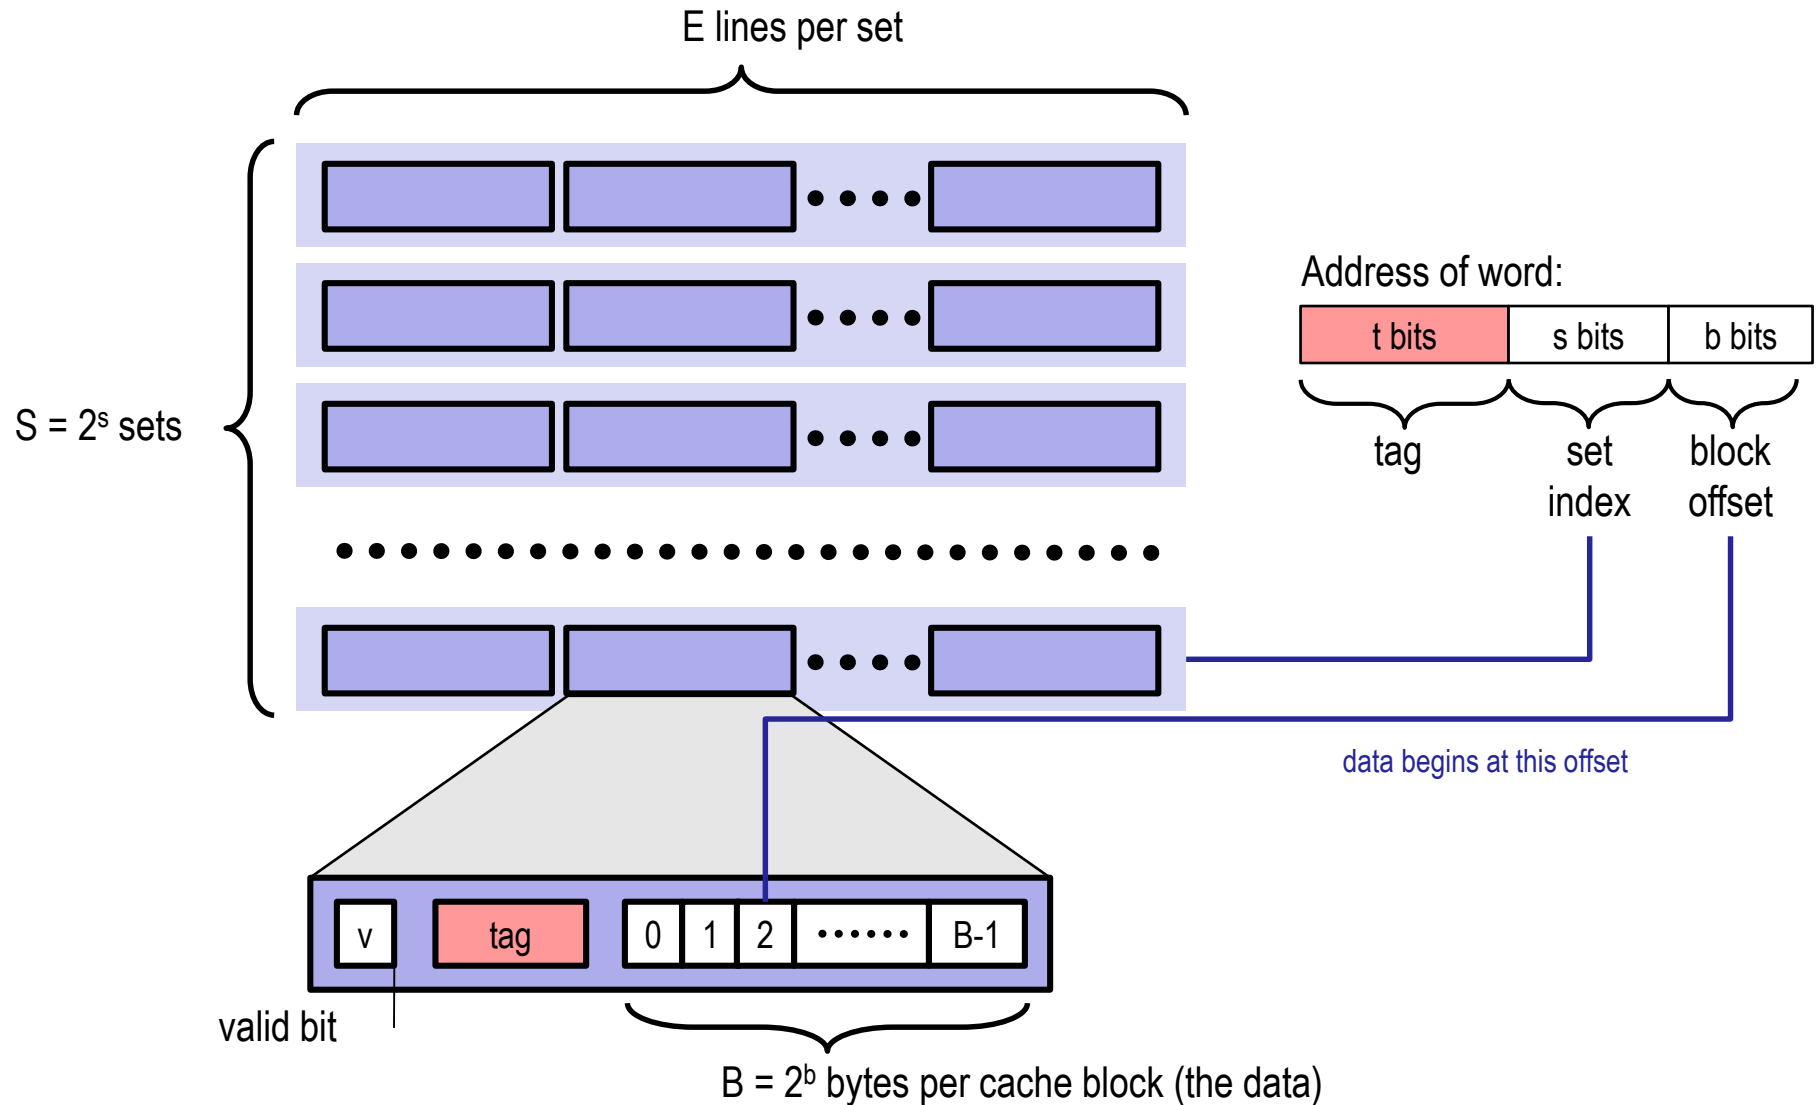


- Block offset: b bits
- Set index: s bits
- Tag Bits: $(\text{Address Size} - b - s)$

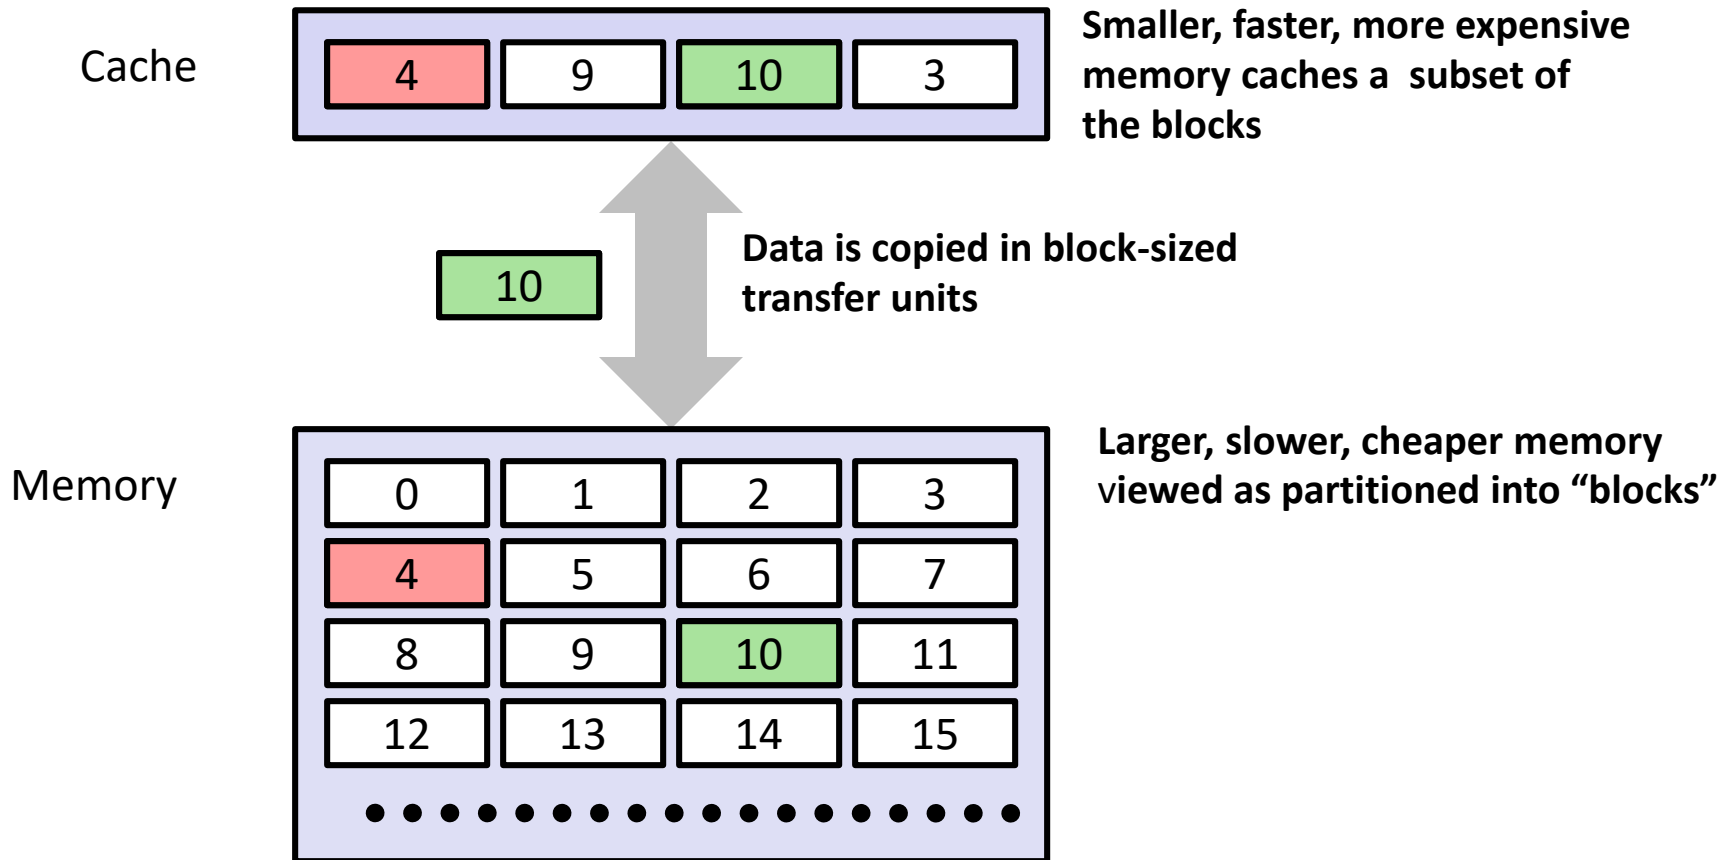
Cache

- A cache is a set of 2^s *cache sets*
- A *cache set* is a set of E *cache lines*
 - E is called associativity
 - If $E=1$, it is called “direct-mapped”
- Each *cache line* stores a block
 - Each block has $B = 2^b$ bytes
- Total Capacity = $S*B*E$

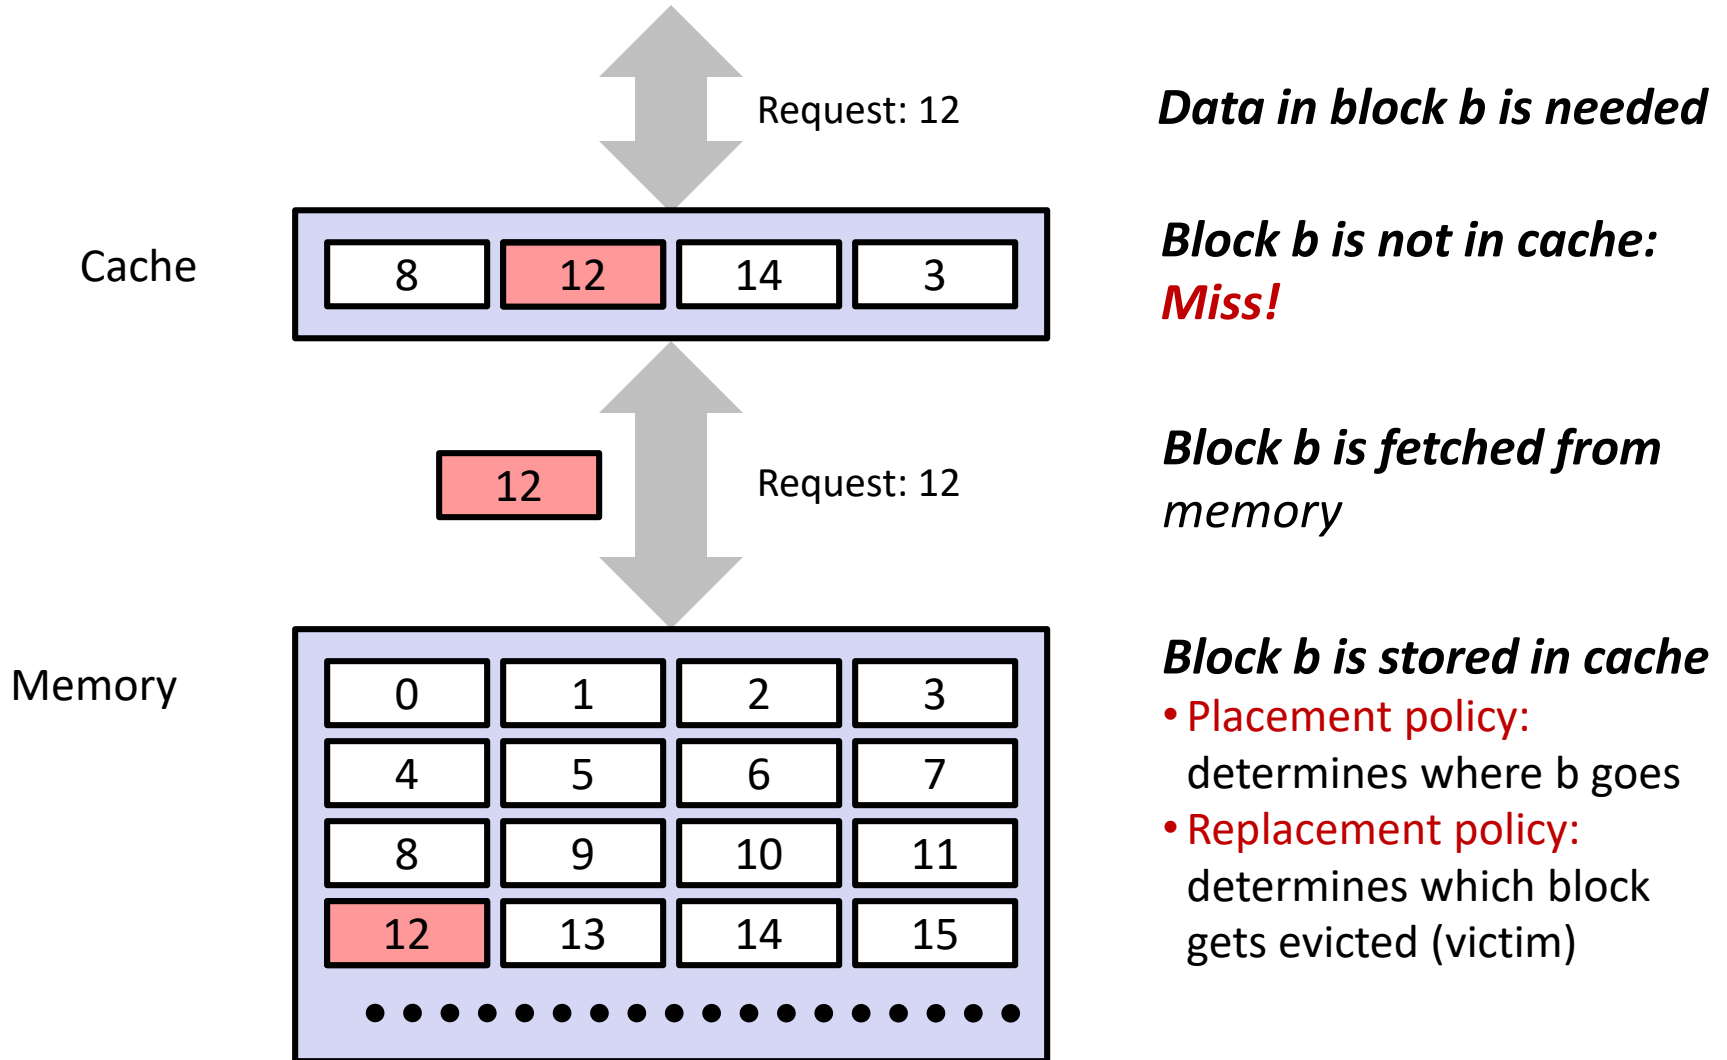
Visual Cache Terminology



General Cache Concepts



General Cache Concepts: Miss



General Caching Concepts:

Types of Cache Misses

■ Cold (compulsory) miss

- The first access to a block has to be a miss

■ Conflict miss

- Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block
 - E.g., Referencing blocks 0, 8, 0, 8, 0, 8, ... would miss every time

■ Capacity miss

- Occurs when the set of active cache blocks (**working set**) is larger than the cache

Cache Simulator

- **A cache simulator is NOT a cache!**
 - Memory contents NOT stored
 - Block offsets are NOT used – the b bits in your address don't matter.
 - Simply **count** hits, misses, and evictions
- **Your cache simulator needs to work for different s , b , e , given at run time.**
- **Use LRU – Least Recently Used replacement policy**
 - Evict the least recently used block from the cache to make room for the next block.
 - Queues ? Time Stamps ?

Cache structure

- **A cache is just 2D array of *cache lines*:**

- `struct cache_line cache[S][E];`
- $S = 2^s$, is the number of sets
- E is associativity

- **Each `cache_line` has:**

- Valid bit
- Tag
- LRU counter (only if you are not using a queue)

getopt

- **getopt() automates parsing elements on the unix command line If function declaration is missing**
 - Typically called in a loop to retrieve arguments
 - Its return value is stored in a local variable
 - When getopt() returns -1, there are no more options
- **To use getopt, your program must include the header file `#include <unistd.h>`**
- **If not running on the shark machines then you will need `#include <getopt.h>`.**
 - Better Advice: Run on Shark Machines !

getopt

- **A switch statement is used on the local variable holding the return value from getopt()**
 - Each command line input case can be taken care of separately
 - “optarg” is an important variable – it will point to the value of the option argument
- **Think about how to handle invalid inputs**
- **For more information,**
 - look at man 3 getopt
 - http://www.gnu.org/software/libc/manual/html_node/Getopt.html

getopt Example

```
int main(int argc, char** argv){
    int opt,x,y;
    /* looping over arguments */
    while(-1 != (opt = getopt(argc, argv, "x:y:"))){
        /* determine which argument it's processing */
        switch(opt) {
            case 'x':
                x = atoi(optarg);
                break;
            case 'y':
                y = atoi(optarg);
                break;
            default:
                printf("wrong argument\n");
                break;
        }
    }
}
```

- Suppose the program executable was called “foo”. Then we would call “./foo -x 1 -y 3” to pass the value 1 to variable x and 3 to y.

fscanf

■ The **fscanf()** function is just like **scanf()** except it can specify a stream to read from (**scanf** always reads from **stdin**)

- parameters:
 - A stream pointer
 - format string with information on how to parse the file
 - the rest are pointers to variables to store the parsed data
- You typically want to use this function in a loop. It returns -1 when it hits EOF or if the data doesn't match the format string

■ **For more information,**

- `man fscanf`
- <http://crasseux.com/books/ctutorial/fscanf.html>

■ **fscanf will be useful in reading lines from the trace files.**

- `L 10,1`
- `M 20,1`

fscanf example

```
FILE * pFile; //pointer to FILE object

pFile = fopen ("tracefile.txt","r"); //open file for reading

char identifier;
unsigned address;
int size;
// Reading lines like " M 20,1" or "L 19,3"

while(fscanf(pFile," %c %x,%d", &identifier, &address, &size)>0)
{
    // Do stuff
}

fclose(pFile); //remember to close file when done
```

Malloc/free

- Use malloc to allocate memory on the heap
- Always free what you malloc, otherwise may get memory leak
 - `some_pointer_you_malloced = malloc(sizeof(int));`
 - `Free(some_pointer_you_malloced);`
- Don't free memory you didn't allocate

Page Replacement Algorithms

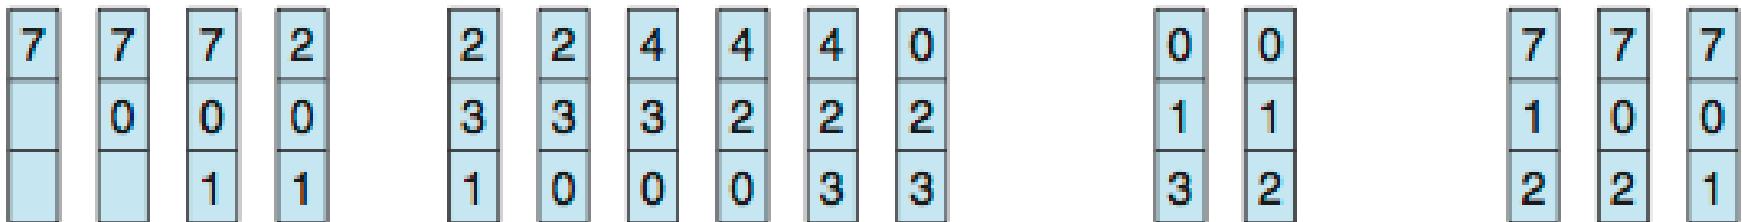
- **When cache is full, a cached data should be replaced with the new data**
 - The new data is referred right now and just used
 - To find a space to hold the new data, choose a data which was previously cached as a victim
 - The victim data is least likely used
- **Algorithms**
 - First-In-First-Out (FIFO) Algorithm
 - Optimal Algorithm
 - Least Recently Used (LRU) Algorithm
- **Implementation of the algorithms in the cachelab**
 - LRU should be implemented by default
 - FIFO and optimal algorithms could be implemented for extra credits

First-In-First-Out (FIFO) Algorithm

- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**
- Cache: Direct Mapped, 3 sets are instantiated
 - Note that the number of sets must be 2^s in your actual cachelab.
 - Page: data block

reference string

7 0 1 2 0 3 0 4 2 3 0 3 0 3 2 1 2 0 1 7 0 1



page frames

- Totally 15 cache misses occurred

Optimal Algorithm

- Replace page(data block) that will not be used for longest period of time
- Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

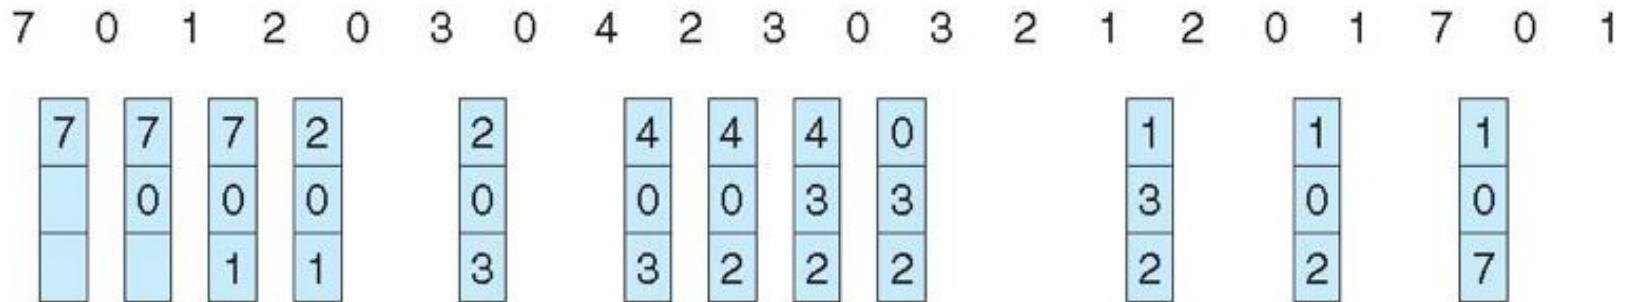


page frames

- Totally 9 cache misses occurred
- Issue: how do you know this?
 - Can't read the future

Least Recently Used (LRU) Algorithm

- **Use past knowledge rather than future**
 - Replace page that has not been used in the most amount of time
- Associate time of last use with each page



- 12 misses – better than FIFO but worse than OPT
- **Generally good algorithm and frequently used**
- **But how to implement?**

LRU Implementation

■ Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- When a page needs to be changed, look at the counters to find smallest value
 - Search through table needed

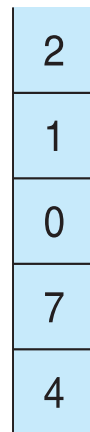
■ Stack implementation

- Keep a stack of page numbers in a double link form:
- Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
- But each update more expensive
- No search for replacement

LRU Implementation Example with Stack

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a



stack
after
b

